



UX RESEARCHER • FOUNDER OF SAGA

Jesse Livingston

I turn scattered signals into one clear story.

A researcher and builder, and a former orchestra conductor.

11 years of UX research: LinkedIn, Netflix, Microsoft, ServiceNow. My throughline is trust: whether users rely on a product, and what makes them stop.

I take fuzzy questions and turn them into rubrics teams can act on. Now building **Saga**, which evaluates how users experience AI.



LinkedIn

NETFLIX

Microsoft

servicenow

[\[email protected\]](#) www.jesselivingston.com

CONTENTS

Selected work

- | | | |
|----|---------------------------------|--|
| 01 | Saga | Evaluating the experience of AI at scale |
| 02 | Microsoft Copilot | Helping users guide the AI experience |
| 03 | LinkedIn Talent Insights | Data backed by transparency and trust |
| 04 | Treehouse | Earning trust from burned car buyers |
-



Evaluating the experience of AI at Scale

Saga, an AI-moderated research platform I built

COMPANY

Saga Insights

ROLE

Founder & CEO

TIMELINE

2025 — Present

METHODS

AI-moderated interviews · rubric design · LLM-as-judge · synthesis · benchmarking

THE PROBLEM

AI is hard to evaluate because it's different for everyone.

A normal feature works the same for everyone, so you can test it with a script. An AI feature answers each person differently, so there's no single path to check.


Thumbs up or down can't tell you why but interviewing every user doesn't scale.

A NORMAL FEATURE



One path you can script


AN AI FEATURE



A new path every time

HOW SAGA WORKS

So I built a platform that catches users in the moment, runs an interview, and scores the result.



SCORING FRAMEWORK


Three dimensions that are critical to AI experiences.

Was it effortless? Do users trust it? Is it solving a real problem they care about? Ease, Trust, and Value are the three pillars.



Saga

317 interviews · live



"I wasn't sure where the number came from, so I double-checked it myself."

TRUST · c-1184 · verified

RESEARCH FOUNDATIONS

A rubric grounded in robust psychometric research.

Each dimension Saga scores is anchored in peer-reviewed work, then tuned for AI-moderated interviews vs. surveys.

SUPR-Q · SAURO · 2015 Factor analysis across 100+ sites shows UX isn't one number: trust loads as its own axis, statistically separate from ease. That's why Saga scores dimensions, not a single grade.

● Ease

SUS Brooke · 1986

Perceived ease is measurable and benchmarkable, separate from task completion.

SEQ Sauro & Dumas · 2009

Ease is best captured in the moment, not in hindsight.

● Trust

Web Credibility Fogg, Stanford · 2003

Trust is a design problem, driven by surface cues.

Trust Repair Kim et al. · 2004

Competence and integrity violations recover differently.

● Value

TAM Davis · 1989

Perceived usefulness drives adoption more than ease.


CPV Zeithaml · 1988

Value is subjective: what I get for what I give.

THE PRODUCT

A living dashboard of insights and scores.

Saga is a fully built product in a real, production environment, built on Postgres, TypeScript, and Anthropic.



Top issue, with the verified quote · spotlights graded on Ease, Trust, Value · what changed since the last run

Helping users guide an AI experience

Microsoft Copilot in PowerPoint

COMPANY

Microsoft (AI & Storytelling)

ROLE

Senior User Researcher, Contract

TIMELINE

Mar — Aug 2024

METHODS

In-depth Interviews & Concept Testing

THE PROBLEM

Microsoft was about to launch a Copilot experience in PPT, but didn't know what users might need from that experience.

PMs had to make launch calls within 3 to 4 weeks, and design was already building, so my work had to shape decisions without slowing them down.

THE CONSTRAINT

3–4 weeks.

Mockups already in flight.

HOW I APPROACHED IT

Two focused studies, scoped to fit the timeline.

I ran 6 one-hour interviews for each study with participants who build work presentations, chosen for a mix of AI experience, company size, and role.

The first study focused on user needs, mental models, and behaviors of using AI to create presentations.

We then incorporated those findings into the designs and launched a second, design-focused study on the prototype.

STUDY DESIGN

WHO

Participants who build work presentations

HOW MANY PARTICIPANTS PER STUDY

6 to 8 interviews

60min

WHAT I STUDIED

1. Foundational needs, mental models, JTBD
2. Design usability

WORKING WITH THE TEAM

Insights in weeks, not months.

I ran the foundational study and the design study in 4 weeks, sharing signals as they came so PMs could decide and designers could adjust inside each two-week cycle.

I made sure every findings was specific and actionable so the PMs could make spec changes and the designers could make UI changes fast.

TWO CYCLES IN FOUR WEEKS

WEEKS 1 — 4

● Research ● Design

CYCLE 1 · WK 1-2

Interviews → insights



Mockups adjust

CYCLE 2 · WK 3-4

Test → refine



Decisions locked

Two full research-and-design loops inside four weeks, not one.

KEY INSIGHT

Participants needed to believe the AI could do what it said it could.

Participants couldn't imagine Copilot making their kind of deck: the weekly update, the progress report. How could Copilot know what their team has been up to?

With that as a blocker, they were hesitant to even try.

The answer? They needed more human touchpoints, more ways to guide the AI experience throughout the process.

HOW MIGHT WE

How might we nudge deck builders to try Copilot for PPT, in spite of their doubt?

RECOMMENDATIONS & IMPACT

We built controls to steer the experience.

Instead of one prompt, users could shape the result. Several of these shipped.

✓ **Start from a file**

Begin from an existing deck or document, not a blank prompt.

✓ **Outline editor**

Edit the outline before Copilot builds the slides.

✓ **Inline AI edits**

Adjust any text or image in place, not by re-prompting.

✓ **Upload a company template**

Keep generated decks on-brand from the first draft.

Data backed by transparency and trust

LinkedIn Talent Insights

COMPANY

LinkedIn

ROLE

Lead UX Researcher

TIMELINE

2019

METHODS

8-10 ninety-min interviews · co-creative workshop · journey mapping ·
concept evaluation

THE SETUP

A new analytics product, and no clear picture of who it was for.

LinkedIn was building a self-serve dashboard of talent data, converting what used to be static reports into something dynamic. But they needed to understand who would be using this, and what specifically they needed.

The team brought me in to find out.

THE STARTING POINT



Years of reports

Shipping, and well used



No picture of the user


Who they were, or what they needed

APPROACH & PERSONAS

Four rounds of research, from foundational to design.

Each round was 8–10 IDIs talking to across talent executives, analysts and data scientists, and recruiters and sourcers. I built proto-personas to give the team one shared picture of who they were designing for.


PROTO-PERSONAS
THE TALENT ECOSYSTEM



JOURNEY MAP

One journey, four stages.

The strategic talent decision-making process we mapped from our interviews.



1 Understand

Read the business strategy and what it implies for talent demand.

2 Analyze

Size internal supply, analyze external supply, and measure demand to gauge hiring difficulty.

3 Plan

Formulate a buy-vs-build-vs-temp plan across quantity, quality, and location.

4 Act

Execute the external recruiting strategy, and pivot if it isn't working.

At every stage the question was the same: **can I defend this number when someone pushes back?**

THE WORKSHOP


I took the personas to ten heads of talent.

I wanted to pressure test our personas, so I ran a co-creative workshop with heads of talent from ten different companies.

They added the context interviews alone could not give me: how these personas and workflows actually show up inside their own organizations.

That grounded the model in more than one company's reality, and showed where it held and where it varied.

TEN HEADS OF TALENT • TEN COMPANIES



The model, grounded

The workshop gave us real examples of our personas in action.

THE TURN

... but then the team got stuck arguing about features.

Like many teams, we ended up getting distracted by feature debates. Should we create a presentation builder? What should notifications look like? What about benchmarking?

Rather than pick a side, I wanted to refocus us on what really mattered. I went back to the data and asked: **what were users raising unprompted, again and again?**

Dashboards?

Notifications?

Presentation builder?

Exports?

Alerts?

Benchmarks?

BACK TO THE DATA Four rounds → one throughline

KEY INSIGHT

It was all about the data, and whether users could defend it.

Participants asked where numbers came from, how they were calculated, and whether they could trust them.

The real fear was looking less credible when an exec or hiring manager pushed back. If they can't explain or defend the data, that doesn't reflect well on them.

*"Where does this data come from?"
"How was this number calculated?"*

TALENT ANALYST, FORTUNE 500

FRAMEWORK & IMPACT

The T.R.A.C. Framework.

I developed this framework to focus all our product decisions and keep this core insight in view.

T

Transparency

"How this is calculated" explainers

R

Relevance

Filtering, folders, the right cut of data

A

Accuracy

Data scientists hired to clean the source

C

Control

Export and adjust to defend the numbers

OUTCOME

\$50M

The fastest product to reach \$50 million in revenue in LinkedIn's history.





Earning trust from burned car buyers

Treehouse EV Marketplace

COMPANY

Treehouse

ROLE

Lead / Sole UX Researcher

TIMELINE

2024

METHODS

7 sixty-min interviews · experience sampling (n=125) · survey (n=250)

THE PROBLEM

An online EV startup that needed strategy research.

The founders had strong instincts but no direct research on foundational user needs. They brought me in to do several studies to understand those needs and shape the strategy.

THREE ASSUMED VALUE PROPS

01 Best prices

02 Easy comparison

03 A concierge to help


HOW I APPROACHED IT

I owned it end to end, with three methods that built on each other.

I set the research strategy and the core questions, then ran the whole thing solo: recruiting, tools, materials, analysis, and the readout.

I recruited for a real spread of EV knowledge and income, and picked tools against budget and speed.

THREE METHODS, IN SEQUENCE



WORKING WITH THE TEAM, AND THE TRADE-OFFS

I scoped each study to what the answer was worth.

I planned bigger and deliberately scaled back once each study was telling me the same thing, because a small startup's time and money are the real constraint.

I worked with the founders as a consultant, translating findings into brand, pricing, and product calls, and re-ran the research after the pivot.

PLANNED VS. ACTUAL

SCOPED TO THE ANSWER

Interviews

~~15~~ → 7



Experience sampling

~~1,000~~ → 125



Survey

~~400~~ → 250



Actual Planned

KEY INSIGHT

The real barrier wasn't price. It was trust.

Almost everyone had been burned by a dealership: shifting prices, pressure, fine print.

1

The Expert

Wants straight, trustworthy detail, an expert path.

2

The Amateur

Wants a guided walkthrough, an assisted path.

"They raised it by ten thousand just to drop it by five. That's not a deal."

EV BUYER, INTERVIEW 04

IMPACT

It reset what the company stood for.

I turned the specific trust-breakers into a list the team could design against.

It shaped the homepage, pricing, and investor story.

THE BRAND, RESET

~~"Efficient car buying"~~

→ **"Trusted car buying"**

Thank you

EMAIL [\[email_protected\]](mailto:[email_protected])

WEB www.jesselivingston.com

LINKEDIN www.linkedin.com/in/jesseliv

